

HW peripherals as software abstractions

Pekka Nikander

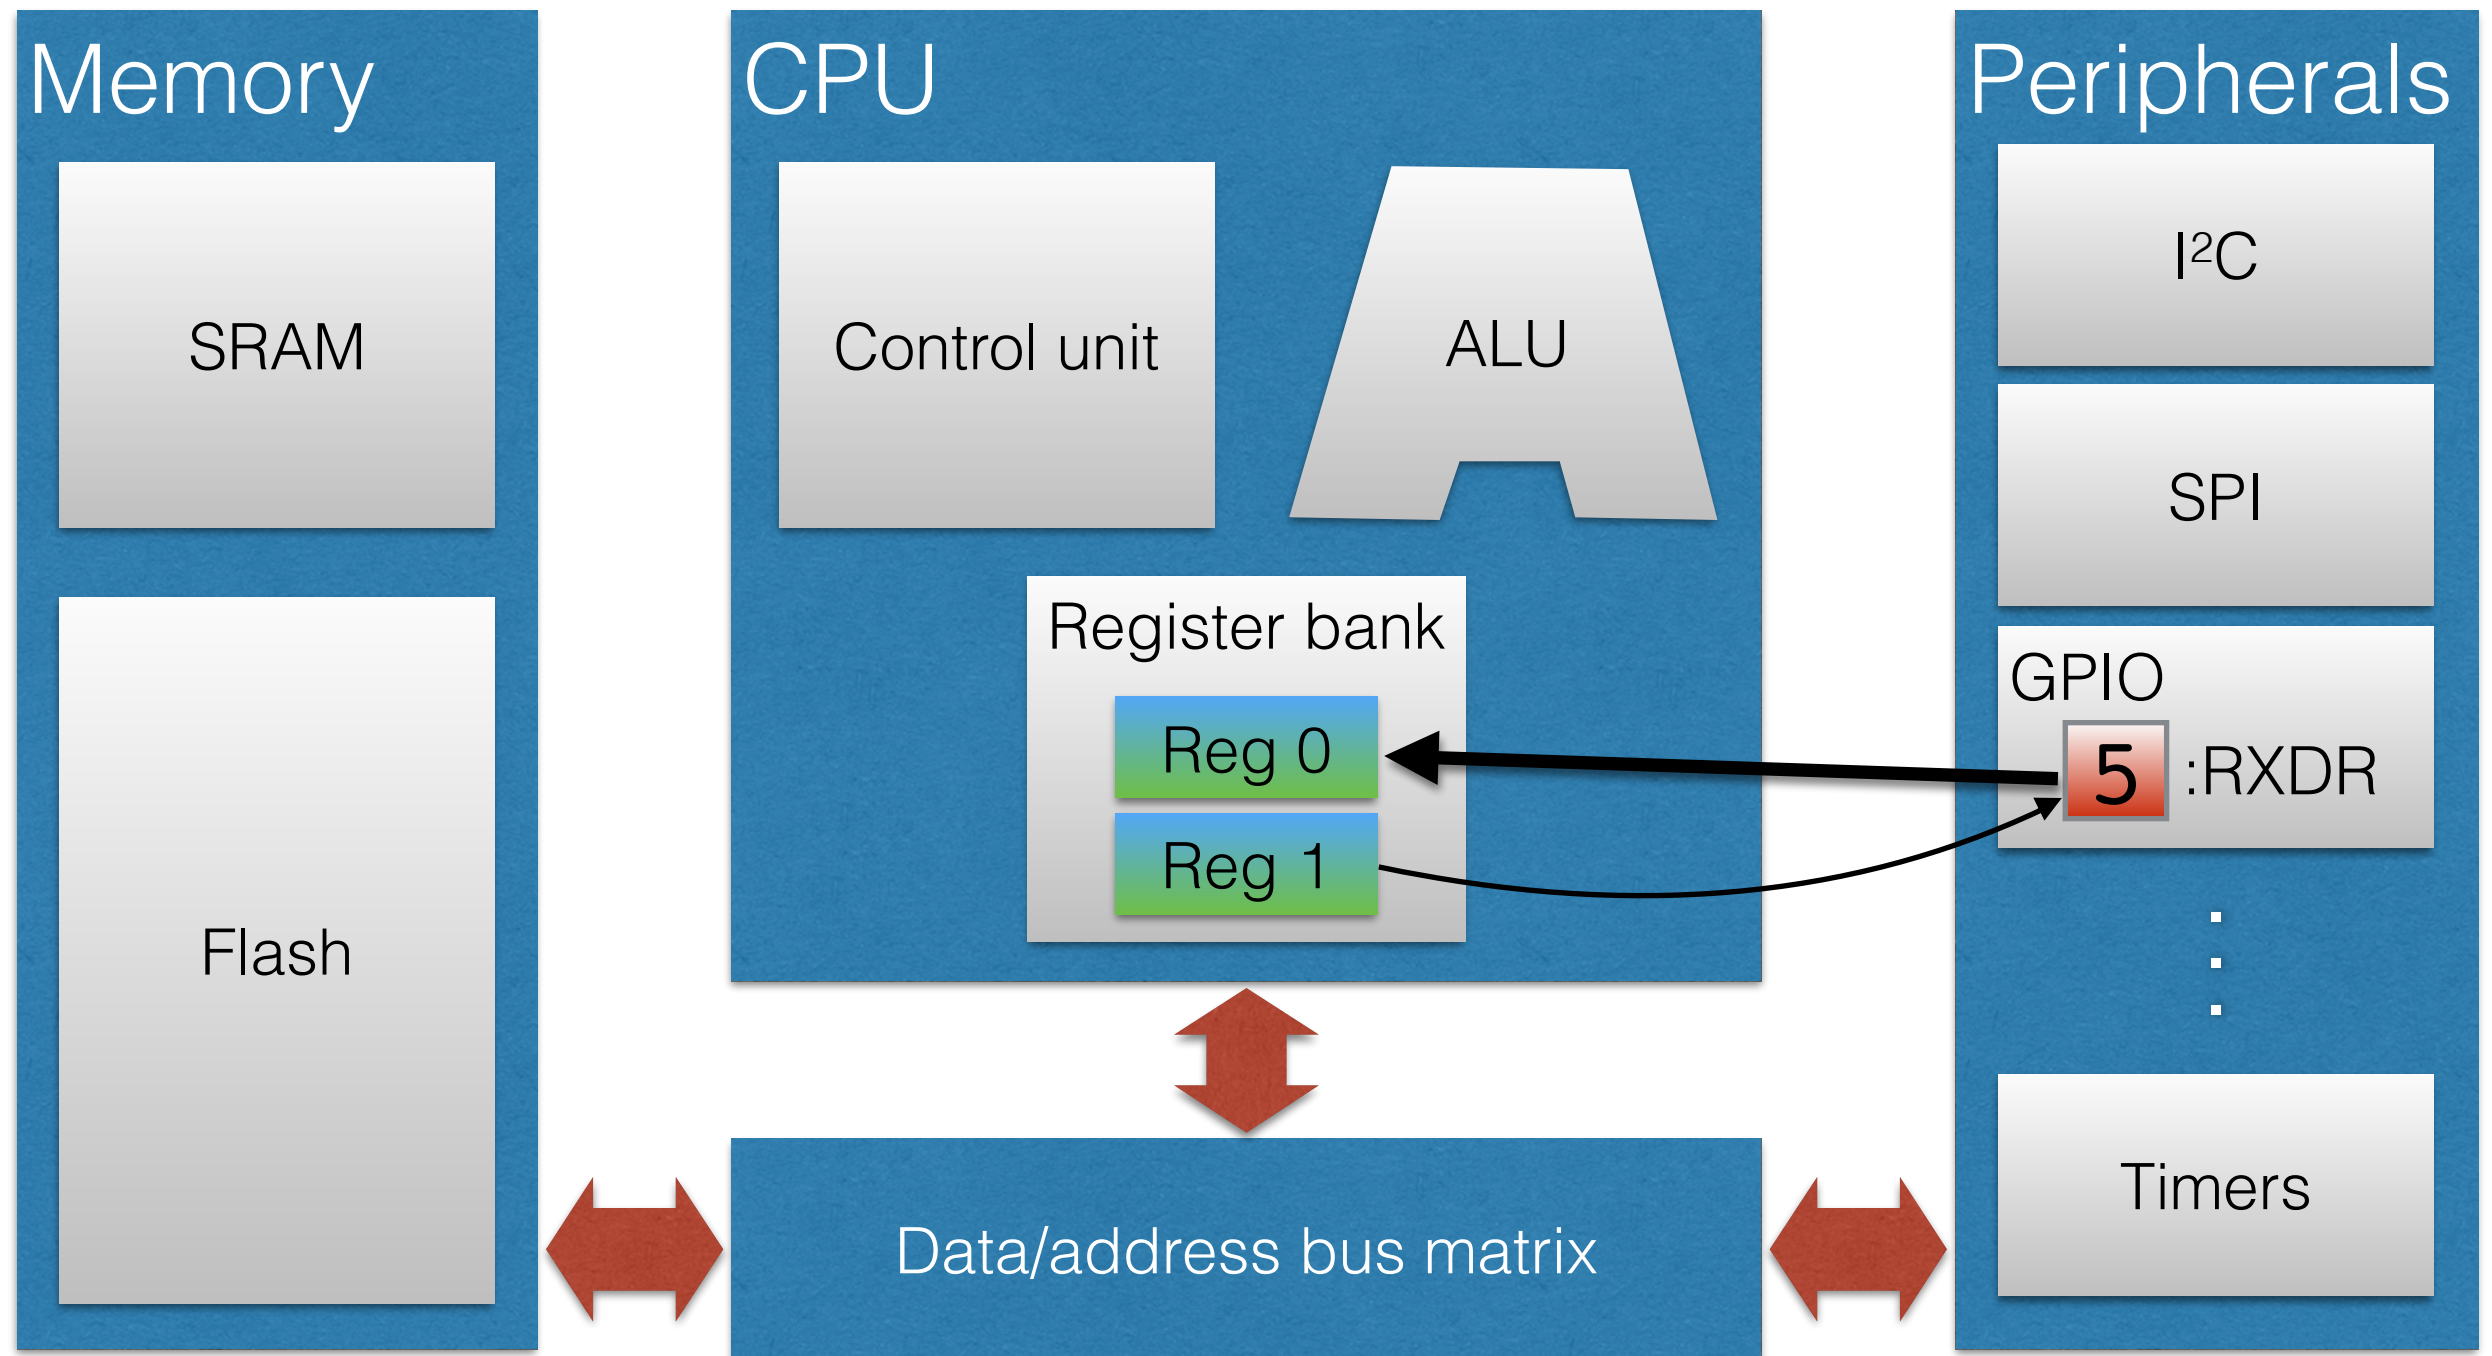
May 2014

Presentation outline

- Recap of memory mapped peripherals
- Basic data type sizes
- Compiler and linker interplay
- Fixed address variables in C / C++
- Meet the optimiser, part I
- Summary



Peripheral registers



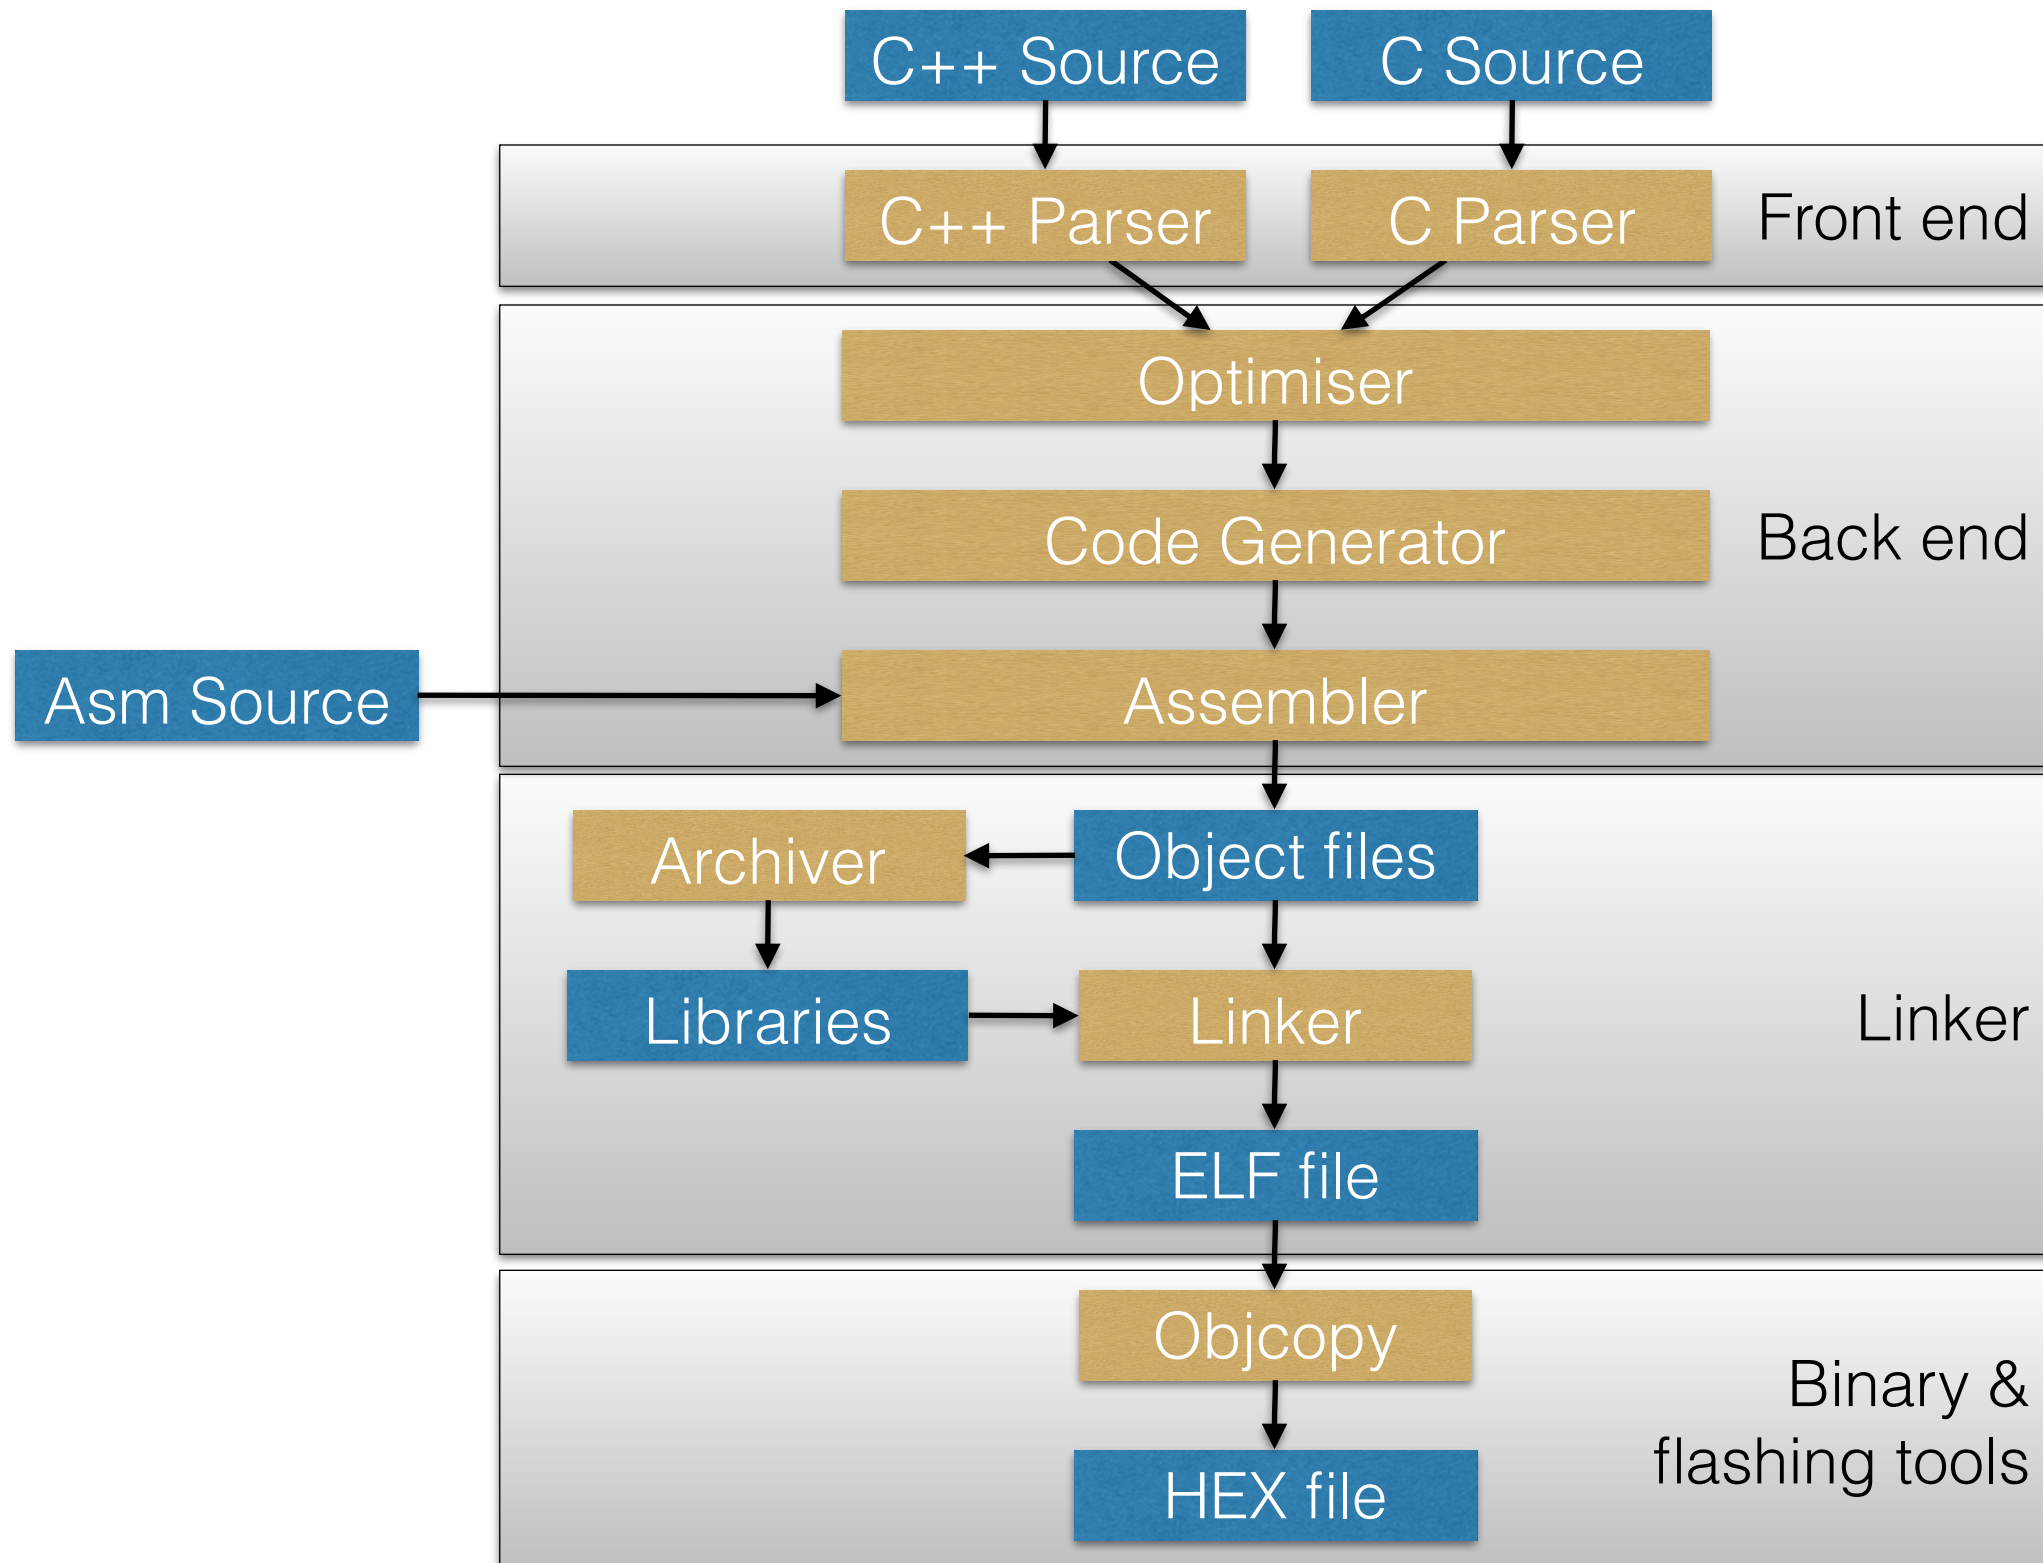
Basic data types

- C was designed to be a *portable* language
 - More modern languages (Java, ...) do it much better
- When doing bare metal, better to be explicit; see the table.
- Avoid unnecessary type casts
 - E.g. prefer unnamed unions to `void` pointers

What	Type
8	<code>uint8_t</code>
16	<code>uint16_t</code>
32	<code>uint32_t</code>
Integer for arithmetic*	<code>int32_t</code>
Generic pointer (to be avoided)	<code>void *</code>

* ARM and thumb generate more efficient code with 32 bit integers

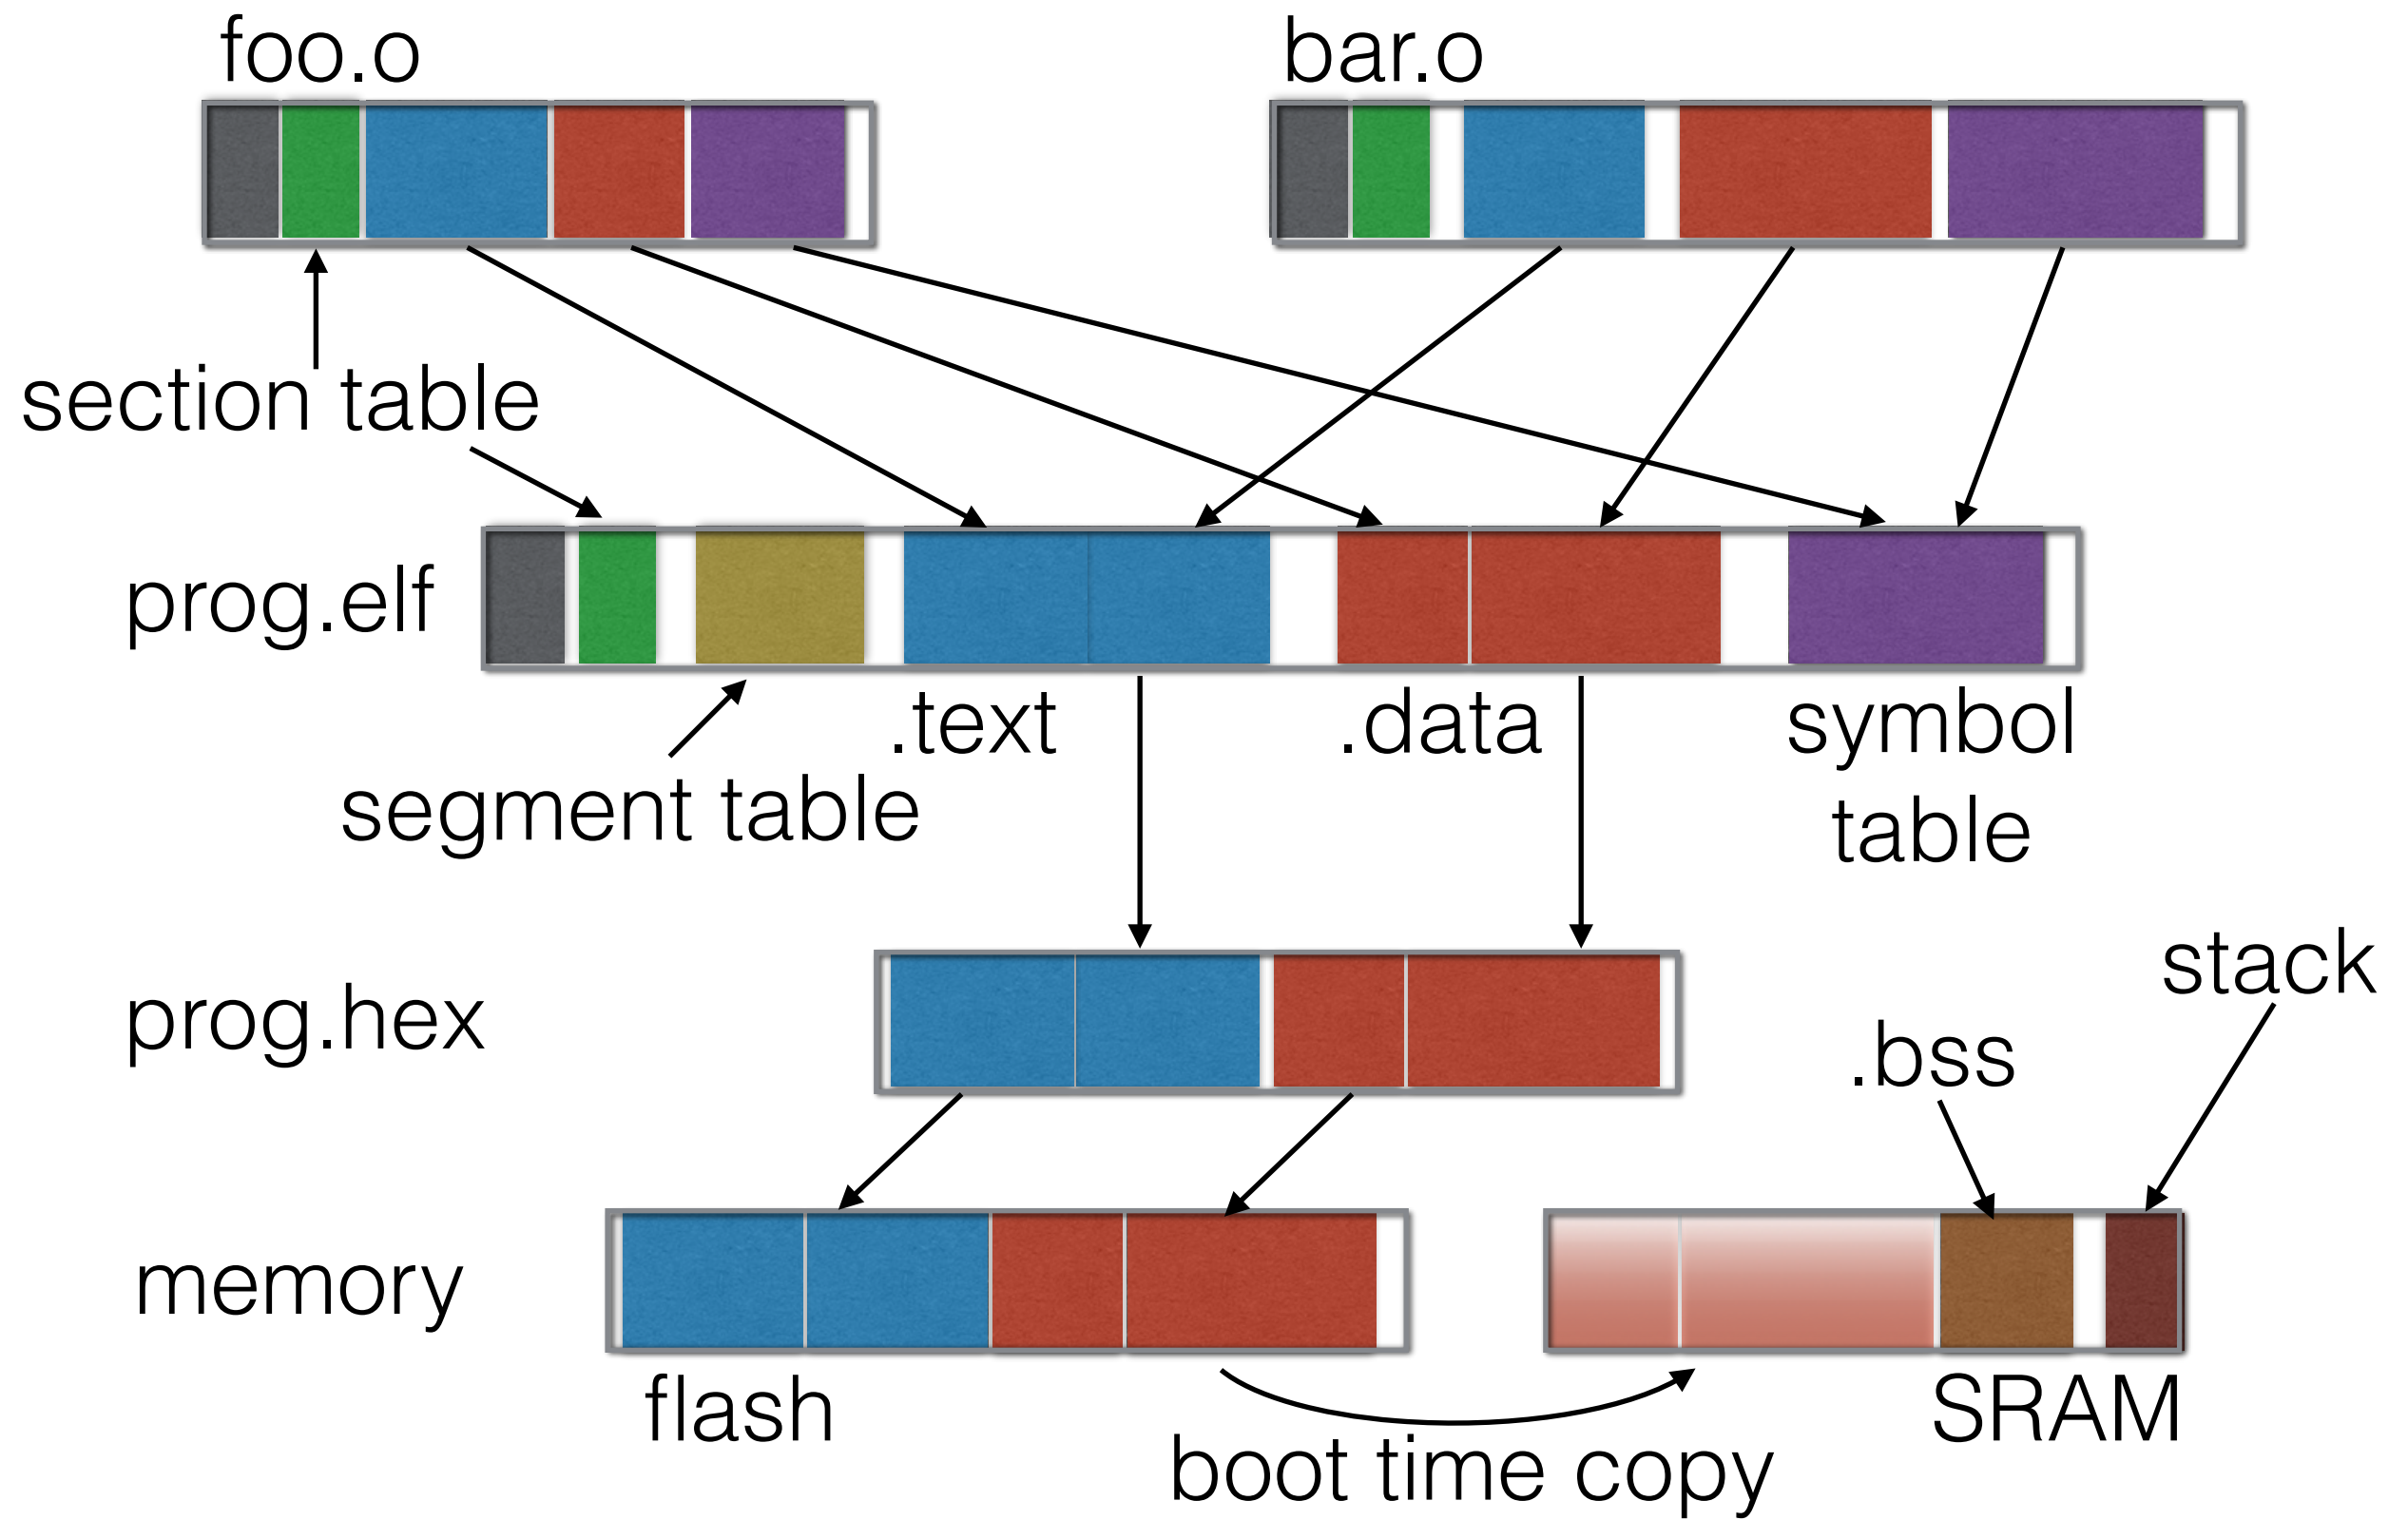
Compiler tool chain



Main components

- Front end: Converts source code to intermediate
- Back end: Analyses, optimises,
generates assembler code
- Assembler: Generates binary object files
- Linker: Combines object files into
an executable
 - Resolves symbols into addresses

Object code anatomy



GNU Linker script

- Defines how the ELF file is constructed
 - Which sections go to which segment
 - Where the segments start in the address space
- Is able to define values (i.e. addresses) for symbols
 - E.g. locate a given variable at a specific address
- Be aware of which linker script is being used


Fixed address variables

- A simple example: a 32-bit integer
 - Address in source code vs. linker script
- A fixed address struct
 - Alignment and padding
- Peripherals in C and C++
 - *POD* in C++ and *standard layout* in C++11

The simple example

- Load I²C read data (RXDR) peripheral register address to the CPU register R1
- Load “memory” from the address in R1 to R0
- Reads the I²C peripheral register RXDR through the data bus matrix to the CPU register R0
- Return the value in R0

```
{  
    volatile int *p;  
    p = 0x40005424;  
    return *p;  
}
```



What's wrong?

- Context completely lost
 - What does `0x40005424` mean?
- Address defined in the C source code
 - A header file would be better
 - The linker script *might* be even better
- (Besides, generates inefficient ARM thumb code)

#1: Use a header file

header.h:

```
#define I2C1_RXDR (volatile int *)0x40005424
```

source.c:

```
{  
    *I2C1_RXDR = 5;  
}
```

#2: Use a struct & define

header.h:

```
struct I2C {  
    ...  
    volatile int RXDR;  
};
```

```
#define I2C1 (struct I2C *const)0x40005400
```

source.c:

```
{  
    I2C1->RXDR = 5;  
}
```

#3: Use a struct & linker

header.h:

```
extern struct I2C {  
    ...  
    volatile int RXDR;  
} I2C1[]; // note the brackets for defining an “array”
```

linker.ld:

```
I2C1 = 0x40005400;
```

source.c:

```
{  
    I2C1->RXDR = 5;  
}
```

Structs & padding

- Be aware of alignment, padding, and endianness
 - Especially with bare metal & communication
- Again, better to be explicit
 - Avoid implicit padding, add explicit fields instead
- Note! Some ARM instructions fail on unaligned data
- Learn to use `__attribute__((...))`

What's wrong here?

```
#define ETHER_ADDR_LEN 6
typedef uint32_t in_addr_t;

struct arp {
    uint16_t    hardware_address_space;
    uint16_t    protocol_address_space;
    uint8_t     hardware_address_length;
    uint8_t     protocol_address_length;
    uint16_t    opcode;
    uint8_t     src_eth_addr[ETHER_ADDR_LEN];
    in_addr_t   src_ip_addr;
    uint8_t     dst_eth_addr[ETHER_ADDR_LEN];
    in_addr_t   dst_ip_addr;
};
```

Fixed

```
#define ETHER_ADDR_LEN 6
typedef uint32_t in_addr_t;

struct arp {
    uint16_t    hardware_address_space;
    uint16_t    protocol_address_space;
    uint8_t     hardware_address_length;
    uint8_t     protocol_address_length;
    uint16_t    opcode;
    uint8_t     src_eth_addr[ETHER_ADDR_LEN];
    in_addr_t   src_ip_addr;
    uint8_t     dst_eth_addr[ETHER_ADDR_LEN];
    in_addr_t   dst_ip_addr;
} __attribute__((packed));
```

C++ PODs

- C++ structs vs. classes
- POD = Plain Old Data
- Making POD structs and classes

C++ structs vs classes

- The *only* difference:
 - `struct` members are `public` by default
 - `class` members are `private` by default
- As a matter of *style*:
 - use the `struct` for something that be a struct in C
 - use the `class` if it uses C++-specific features
- <http://stackoverflow.com/questions/7762085/difference-between-a-struct-and-a-class>

Plain Old Data (POD)

- Type (incl. classes) where the compiler guarantees that there is no “magic”, for example
 - no hidden pointers to vtables
 - no offsets that get applied to the address cast
 - no constructors or destructors
- Roughly speaking, a type is a POD when the only things in it are built-in types and their combinations

Why to consider PODs?

- Define exact binary-level memory layout
- Allow structs (or unions) to be coated with syntactic sugar
 - Define member functions
 - Define static data
- Allows going to C++ and still “think in C”

C++11 changes

- Goal: a POD in C++11 gives you the same memory layout as a struct compiled in C
 - A POD is both trivial and standard layout
- Standard layout is often enough, but be careful if you decide to use non-trivial features
 - Possible to use (limited) inheritance, constructors, protected or private, etc
- <http://stackoverflow.com/questions/4178175/what-are-aggregates-and-pods-and-how-why-are-they-special>

Rules of thumb

- Definitely OK to define a `struct` or `union` as in C
- The following are not OK
 - Virtual functions
 - Multiple inheritance
 - Constructors that are not `constexpr`

A C++11 example

header.h:

```
class I2C {
    volatile int RXDR;
public:
    size_t read(uint_t *buf, size_t bufsize);
} I2C1[]; // note the brackets for defining an "array"

inline size_t I2C::read(uint_t *b, size_t *b) {
    while (...) {
        ...
        b[c] = RXDR;
    }
}
```

linker.ld:

```
I2C1 = 0x40005400;
```

Meet the optimiser

Part I



A few common tricks & traps

- Inlining
- Constant propagation
- LLVM and undefined behaviour
- Do volatilise

Inlining

- Expand function at where it is called
 - Avoids function call overhead (not so important)
 - Allows other optimisations to be applied
 - Especially constant propagation
- If a function is called only once, trivially useful
- Otherwise you need to estimate cost vs benefit

Constant propagation

- Evaluate constant expressions at compile time
 - C++11 has even the `constexpr` keyword
- Propagate the newly computed constants, i.e. evaluate more values with their constant values
- Sometimes lead to significant code size reduction
- Some parts of ELL-i runtime are carefully designed to become optimised due to constant propagation

Example: inlining & constant propagation

```
static inline
void digitalWrite(pin_t pin, uint32_t val) {
    if (val)
        GPIOPIN[pin].gpio_port->BSRR = GPIOPIN[pin].gpio_mask;
    else
        GPIOPIN[pin].gpio_port->BRR  = GPIOPIN[pin].gpio_mask;
}

...
digitalWrite(13, 1);
...
```


Step #1: Inlining

```
static inline  
void digitalWrite(pin_t pin, uint32_t val) {  
    if (val)  
        GPIOPIN[pin].gpio_port->BSRR = GPIOPIN[pin].gpio_mask;  
    else  
        GPIOPIN[pin].gpio_port->BRR = GPIOPIN[pin].gpio_mask;  
}
```

```
...  
digitalWrite(13, 1);  
if (HIGH)  
    GPIOPIN[13].gpio_port->BSRR = GPIOPIN[13].gpio_mask;  
else  
    GPIOPIN[13].gpio_port->BRR = GPIOPIN[13].gpio_mask;  
...
```

Step #2: *if* optimisation

```
...  
digitalWrite(13, 1);  
if (1)  
    GPIOPIN[13].gpio_port->BSRR = GPIOPIN[13].gpio_mask;  
else  
    GPIOPIN[13].gpio_port->BRR = GPIOPIN[13].gpio_mask;  
...
```

Step #3: Constant propagation

```
static const struct {} GPIOPIN[] = {  
    ...  
    { .gpio_port = GPIOA, .gpio_mask = 0x02000 },  
    ...  
};  
  
...  
digitalWrite(13, 1);  
if (1)  
    GPIOPIN[13].gpio_port->BSRR = GPIOPIN[13].gpio_mask;  
else  
    GPIOPIN[13].gpio_port->BRR = GPIOPIN[13].gpio_mask;  
...
```

Step #3: Constant propagation

```
static const struct {} GPIOPIN[] = {  
    ...  
    { .gpio_port = GPIOA, .gpio_mask = 0x2000 }, // pin 13  
    ...  
};  
  
...  
digitalWrite(13, 1);  
if (1)  
    GPIOPIN[13].gpioGPIOA->BSRR = 0x2000N[13].gpio_mask;  
else  
    GPIOPIN[13].gpio_port->BRR = GPIOPIN[13].gpio_mask;  
...  

```

Example: End result

...

```
GPIOA->BSRR = 0x2000;
```

...

Undefined behaviour

- Designers of C wanted it to be an extremely efficient
- Therefore some seemingly reasonable things in C have *undefined* behaviour
 - And it really is *completely* undefined
- LLVM could utilise undefined less than it could ...
 - ... but it still produces surprises in many cases
- <http://blog.llvm.org/2011/05/what-every-c-programmer-should-know.html>

An LLVM example

- Consider the following code
 - It checks for the `null` pointer, doesn't it?

```
void contains_null_check(int *P) {  
    int dead = *P;  
    if (P == 0)  
        return;  
    *P = 4;  
}
```


An LLVM example

- But what if the compiler happens to run
 - "Redundant Null Check Elimination" first and
 - "Dead Code Elimination" afterwards

```
void contains_null_check(int *P) {  
    int dead = *P;  
    if (P == 0)  
        return;  
    *P = 4;  
}
```

An LLVM example

- `dead = *P` dereferences `P`
 - Remember, dereferencing `NULL` is *undefined*
- Hence, `P` cannot be `NULL`

```
void contains_null_check(int *P) {  
    int dead = *P;  
    if (false) // P cannot be NULL  
        return;  
    *P = 4;  
}
```

An LLVM example

- Next eliminate dead code

```
void contains_null_check(int *P) {  
    int dead = *P;  
    if (false)  
    return;  
    *P = 4;  
}
```

Once again `volatile`

- LLVM (and GCC) *will* optimise your writes away

```
for (int i = 0; i < 100000; i++)  
    ;
```
- Be sure to declare `volatile` if you don't want that
- Remember
 - LLVM and GCC both try to be helpful
 - But sometimes they aren't, as they try to generate efficient code, and can't read your intentions

Summary

- Feel free to declare peripherals as structs or PODs
 - Consider using the linker to define the address...
- Heavy inlining, with constant propagation, allows you to write high-level code that optimises well
- Beware of *undefined* behaviour, lest thou regret